

Optimization

2D Optimization

Sachin Shanbhag

Department of Scientific Computing
Florida State University,
Tallahassee, FL 32306.



References

- ▶ S. Chapra and R. Canale, Numerical Methods for Engineers
- ▶ A. Greenbaum and T. Chartier, Numerical Methods: Design, Analysis, and Computer Implementation of Algorithms
- ▶ Carnahan and Wilkes, Applied Numerical Methods, University of Michigan, Class Notes, 1996.
- ▶ Burden and Faires, Numerical Analysis, 1993
- ▶ Pal, Numerical Analysis for Scientists and Engineers, 2007.
- ▶ M. Heath, Scientific Computing: An Introductory Survey
- ▶ wikipedia.org

Contents


- ▶ Basics and Background
- ▶ Optimization: Single variable
 - ▶ Golden Section
 - ▶ Quadratic Interpolation
 - ▶ Newton's Method
- ▶ Multidimensional Optimization*
 - ▶ Steepest Descent
 - ▶ Newton's Method
 - ▶ BFGS Method: Quasi-Newton
- ▶ Miscellaneous

*In this class “multi” = 2, although we will sometimes use more general language

Multidimensional functions


- Instead of just $f(x)$, we will now consider finding the *minima* of functions $f(\mathbf{x}) = f(x_1, x_2, \dots, x_n)$
- Example: Consider the 2D function

$$f(x_1, x_2) = 0.5x_1^2 + 2.5x_2^2$$



2D Function

- We can also construct a contour plot



- An important concept in multidimensional optimization is the gradient of $f(\mathbf{x})$. For a 2D function such as the one here:

$$\nabla f(\mathbf{x}) = \frac{\partial f}{\partial x_1} \mathbf{e}_1 + \frac{\partial f}{\partial x_2} \mathbf{e}_2$$

Gradient

- ▶ The gradient generalizes the concept of derivative to multiple dimensions
- ▶ Note that it is a vector, and has a “direction” in addition to a magnitude. This direction is important in optimization.
- ▶ We can also write it as a vector


$$\nabla f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \end{bmatrix}, \quad \text{assuming } \mathbf{e}_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \mathbf{e}_2 = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

- ▶ We can evaluate the gradient of this particular function:

$$\nabla f(\mathbf{x}) = \begin{bmatrix} x_1 \\ 5x_2 \end{bmatrix} = x_1 \mathbf{e}_1 + 5x_2 \mathbf{e}_2$$

Gradient

- I plotted the direction of the gradient vector at a few different points on the contour map



- The three points and the corresponding normalized gradients are $(\mathbf{x}, \nabla f(\mathbf{x}) / \|\nabla f(\mathbf{x})\|)$

$$\left(\begin{bmatrix} 0.0 \\ 2.0 \end{bmatrix}, \begin{bmatrix} 0.0 \\ 1.0 \end{bmatrix} \right), \quad \left(\begin{bmatrix} 3.0 \\ 1.5 \end{bmatrix}, \begin{bmatrix} 0.37 \\ 0.93 \end{bmatrix} \right), \quad \left(\begin{bmatrix} -3.8 \\ 1.0 \end{bmatrix}, \begin{bmatrix} -0.61 \\ -0.80 \end{bmatrix} \right)$$

Gradient

- ▶ Note that the gradient is perpendicular to contour lines
- ▶ The direction of ∇f tells us which way to travel in to gain elevation as quickly as possible
- ▶ The magnitude of ∇f tell us how much we gain by travelling in that direction
- ▶ This is similar to derivative of a single variable $f(x)$ where df/dx measures the “rate” at which $f(x)$ changes with x .
- ▶ Next we are going to consider a method called “steepest descent” to find the *minima* of a function $f(\mathbf{x})$ by travelling in the direction of $-\nabla f(\mathbf{x})$
- ▶ There is completely analogous method called “steepest ascent” to find the maxima by travelling in the direction of $\nabla f(\mathbf{x})$

Steepest Descent: Algorithm

1. $k = 0$; $\mathbf{x}_0 =$ initial guess
2. Compute the -ve gradient $\mathbf{s}_k = -\nabla f(\mathbf{x}_k)$
3. Choose α_k to minimize $f(\mathbf{x}_k + \alpha \mathbf{s}_k)$. Note that this is a 1D problem, for which we have devised methods before. This is called a *line* search.
4. Update the solution: $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{s}_k$
5. Set $k = k + 1$, and go back to step 2, and repeat until convergence.

Example

Problem: Use steepest descent to minimize the function[†]

$$f(\mathbf{x}) = 0.5x_1^2 + 2.5x_2^2,$$

where $\mathbf{x} = [x_1, x_2]^T$, and whose gradient is:

$$\nabla f(\mathbf{x}) = \begin{bmatrix} x_1 \\ 5x_2 \end{bmatrix}$$

Start with an initial guess of $\mathbf{x}_0 = [5, 1]^T$.

Solution:

$$\mathbf{x}_0 = [5, 1]^T, \quad \mathbf{s}_0 = -\nabla f(\mathbf{x}_0) = -[5, 5]^T$$

Therefore,

$$f(\mathbf{x}_0 + \alpha \mathbf{s}_0) = f\left(\begin{bmatrix} 5 \\ 1 \end{bmatrix} - \alpha \begin{bmatrix} 5 \\ 5 \end{bmatrix}\right) = f\left(\begin{bmatrix} 5 - 5\alpha \\ 1 - 5\alpha \end{bmatrix}\right)$$

[†]Problem from Heath, chapter 6.

Example

Thus,

$$\begin{aligned}f(\mathbf{x}_0 + \alpha \mathbf{s}_0) &= 0.5(5 - 5\alpha)^2 + 2.5(1 - 5\alpha)^2 \\&= 75\alpha^2 - 50\alpha + 15\end{aligned}$$

One can easily find the minima of this function by taking the derivative with respect to α which gives us $150\alpha - 50 = 0$, or $\alpha_0 = 1/3$.

Thus,

$$\mathbf{x}_1 = \mathbf{x}_0 + (1/3)\mathbf{s}_0 = \begin{bmatrix} 3.333 \\ -0.667 \end{bmatrix}$$

We can now repeat the process until we are happy!

Or we can write the following matlab code.

Steepest Descent: Matlab Program

```
% Steepest Descent Demo: use [x f n] = steepDescent([5;1], 1e-3)

function [xopt fopt nopt] = steepDescent(x0, tol)

    x = x0;
    k = 0;

    while(norm(gradf(x)) > tol)                                % Need to make gradf = 0

        s      = -gradf(x);
        falpha = @(alpha) f(x + alpha*s);
        alpha  = fminsearch(falpha,0.1);
        x      = x + alpha * s;
        k      = k + 1;

    end

    xopt = x;    fopt = f(x); nopt = k;

end


function Z = f(x)
    Z = 0.5*x(1)^2 + 2.5*x(2)^2;
end

function Z = gradf(x)
    Z = [x(1);5*x(2)];
end
```

Example

From Heath pg 278:

k	x_k^T		$f(x_k)$	$\nabla f(x_k)^T$	
0	5.000	1.000	15.000	5.000	5.000
1	3.333	-0.667	6.667	3.333	-3.333
2	2.222	0.444	2.963	2.222	2.222
3	1.481	-0.296	1.317	1.481	-1.481
4	0.988	0.198	0.585	0.988	0.988
5	0.658	-0.132	0.260	0.658	-0.658
6	0.439	0.088	0.116	0.439	0.439
7	0.293	-0.059	0.051	0.293	-0.293
8	0.195	0.039	0.023	0.195	0.195
9	0.130	-0.026	0.010	0.130	-0.130



What is the geometrical interpretation?

Hessian

- ▶ The Hessian of a multidimensional scalar function $f(\mathbf{x})$ is given by the symmetric square matrix

$$\mathbf{H}_f(\mathbf{x}) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

- ▶ Just as the gradient generalizes df/dx , the Hessian generalizes d^2f/dx^2 .

Hessian

- In 2D, the Hessian is:

$$\mathbf{H}_f(\mathbf{x}) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} \end{bmatrix}$$

- Recall how $d^2 f(x^*)/dx^2$ told us whether x^* (obtained by solving $df/dx = 0$) was a maxima, minima or a saddle point
- The Hessian does the same job. If \mathbf{x}^* is a solution to $\nabla f(\mathbf{x}) = \mathbf{0}$, then if $\mathbf{H}_f(\mathbf{x}^*)$ is
 - +ve definite $\implies x^*$ is a minima
 - ve definite $\implies x^*$ is a maxima
 - indefinite $\implies x^*$ is a saddle point

Newton's Method

- ▶ Recall 1D Newton's Method for optimization:

$$x_{i+1} = x_i - \frac{f'(x_i)}{f''(x_i)}$$

- ▶ We could rewrite this expression as:

$$(x_{i+1} - x_i)f''(x_i) = -f'(x_i)$$

- ▶ We are going to generalize this method for multiple dimensions
- ▶ It is useful to visualize Newton's method as a quadratic approximation to a Taylor's series

Newton's Method

- ▶ That is consider

$$f(x + s) \approx f(x) + \frac{df}{dx}s + \frac{1}{2} \frac{d^2f}{dx^2}s^2.$$

- ▶ We can think of the RHS as a quadratic function in s which can be minimized

$$\frac{df(x + s)}{ds} = 0 \implies 0 + \frac{df}{dx} + \frac{d^2f}{dx^2}s = 0$$

- ▶ Leading to

$$s \frac{d^2f}{dx^2} = -\frac{df}{dx}$$

which is the same as Newton's method for optimization

Newton: Multidimensional case

- We can repeat the Taylor series expansion for $f(\mathbf{x})$.

$$f(\mathbf{x} + \mathbf{s}) \approx f(\mathbf{x}) + \nabla f(\mathbf{x})^T \mathbf{s} + \frac{1}{2} \mathbf{s}^T \mathbf{H}_f(\mathbf{x}) \mathbf{s}$$

Note that for 1D this collapses to the previous expression.

- Taking the derivative, leads us to:

$$\mathbf{H}_f(\mathbf{x}) \mathbf{s} = -\nabla f(\mathbf{x})$$

Compare with the 1D case:

$$s \frac{d^2 f}{dx^2} = -\frac{df}{dx}$$

- This allows us to write an algorithm for Newton's method

Newton's Method: Algorithm

1. $k = 0$; $\mathbf{x}_0 =$ initial guess
2. Compute the gradient $\nabla f(\mathbf{x}_k)$ and the Hessian $\mathbf{H}_f(\mathbf{x}_k)$
3. Solve $\mathbf{H}_f(\mathbf{x}_k)\mathbf{s}_k = -\nabla f(\mathbf{x}_k)$ for \mathbf{s}_k
4. Update the solution: $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{s}_k$
5. Set $k = k + 1$, and go back to step 2, and repeat until convergence.

Example

Problem: Solve the previous example again, this time using Newton's method

$$f(\mathbf{x}) = 0.5x_1^2 + 2.5x_2^2,$$

Solution:

The gradient and Hessian are given by:

$$\nabla f(\mathbf{x}) = \begin{bmatrix} x_1 \\ 5x_2 \end{bmatrix}, \quad \mathbf{H}_f(\mathbf{x}_k) = \begin{bmatrix} 1 & 0 \\ 0 & 5 \end{bmatrix}$$

We solve the system (with $\mathbf{x}_0 = [5, 1]^T$)

$$\begin{bmatrix} 1 & 0 \\ 0 & 5 \end{bmatrix} \mathbf{s}_0 = - \begin{bmatrix} 5 \\ 5 \end{bmatrix} \implies \mathbf{s}_0 = \begin{bmatrix} -5 \\ -1 \end{bmatrix}$$

Example

- ▶ This implies

$$\mathbf{x}_1 = \mathbf{x}_0 + \mathbf{s}_0 = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

which is the true solution

- ▶ It is not surprising that Newton's method converged in 1 iteration since the $f(\mathbf{x})$ was quadratic
- ▶ In general the convergence rate is quadratic, but the method can veer off unless we start close enough to the solution
- ▶ Note: No line search required, but we had to determine a Hessian matrix and solve a linear system at each iteration
- ▶ In damped Newton methods, a line search is added to make the method more robust.

Quasi-Newton Methods

- ▶ Newton's method converges rapidly once you are close to the solution. But it doesn't come cheap.
- ▶ For a n -dimensional problem, each iteration requires $\mathcal{O}(n^2)$ function evaluations to form the gradient and the Hessian, and $\mathcal{O}(n^3)$ operations to solve the linear system.
- ▶ To reduce overhead, quasi-Newton methods have been developed which seek to replace the step:

$$\mathbf{H}_f(\mathbf{x}_k)(\mathbf{x}_{k+1} - \mathbf{x}_k) = -\nabla f(\mathbf{x}_k)$$

with

$$\mathbf{B}_k(\mathbf{x}_{k+1} - \mathbf{x}_k) = -\alpha_k \nabla f(\mathbf{x}_k)$$

where \mathbf{B} is an approximation to the Hessian matrix, and may be obtained by secant updating and α_k is a line-search parameter.

BFGS Method

- ▶ A popular secant updating method named after its co-inventors: Broyden, Fletcher, Goldfarb and Shanno.
- ▶ Initially set $\mathbf{B}_0 = \mathbf{I}$, which means the first step is in the negative gradient direction (like steepest descent).
- ▶ Unlike Newton's method, the second derivatives (Hessian) do not have to be pre-computed.
- ▶ It is built up over time.
- ▶ Convergence is superlinear.
- ▶ We consider a simple algorithm (better implementations update a factorization of the matrix \mathbf{B} rather than the matrix itself)

BFGS Algorithm

1. $k = 0$; $\mathbf{x}_0 =$ initial guess
2. Set $\mathbf{B}_0 = \mathbf{I}$ as the initial Hessian approximation
3. Compute the gradient $\nabla f(\mathbf{x}_k)$
4. Solve $\mathbf{B}_k \mathbf{s}_k = -\nabla f(\mathbf{x}_k)$ for \mathbf{s}_k
5. Update the solution: $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{s}_k$
6. Set $\mathbf{y}_k = \nabla f(\mathbf{x}_{k+1}) - \nabla f(\mathbf{x}_k)$
7. Update the Hessian

$$\mathbf{B}_{k+1} = \mathbf{B}_k + \frac{\mathbf{y}_k \mathbf{y}_k^T}{\mathbf{y}_k^T \mathbf{s}_k} - \frac{\mathbf{B}_k \mathbf{s}_k \mathbf{s}_k^T \mathbf{B}_k}{\mathbf{s}_k^T \mathbf{B}_k \mathbf{s}_k}$$

8. Set $k = k + 1$, and go back to step 4, and repeat until convergence.

Example

Problem: Solve the previous example again, this time using BFGS method

$$f(\mathbf{x}) = 0.5x_1^2 + 2.5x_2^2,$$

Solution:

The gradient and approximate Hessian are given by:

$$\nabla f(\mathbf{x}) = \begin{bmatrix} x_1 \\ 5x_2 \end{bmatrix}, \quad \mathbf{B}_0 = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

We solve the system (with $\mathbf{x}_0 = [5, 1]^T$).

The first step is simply: $\mathbf{Is}_0 = -\nabla f(\mathbf{x}_0) = -[5, 5]^T$

$$\implies \mathbf{x}_1 = [5, 1]^T + [-5, -5]^T = [0, -4]^T$$

We can update the approximate Hessian to

$$\mathbf{B}_1 = \begin{bmatrix} 0.667 & 0.333 \\ 0.333 & 0.667 \end{bmatrix}$$

Example

We can continue to get the sequence:

k	x_1	x_2	$f(\mathbf{x})$
0	5.0000	1.0000	15.0000
1	0.0000	-4.0000	40.0000
2	-2.2222	0.4444	2.9630
3	0.8163	0.0816	0.3499
4	-0.0092	-0.0153	0.0006
5	-0.0005	0.0009	0.0000

using the following Matlab code:

Matlab Code

```
% BFGS Demo: use as [x f n] = bfgs([5;1], 1e-3);

function [xopt fopt nopt] = bfgs(x0, tol)

    x = x0;
    n = length(x);
    B = eye(n);
    k = 0;

    while(norm(gradf(x)) > tol) % Need to make gradf ~ 0
        df = gradf(x);
        s = B\(-df);
        x = x + s;
        y = gradf(x) - df;
        B = B + (y*y')/(y'*s) - (B*s*s'*B)/(s'*B*s);
        k = k + 1;
    end

    xopt = x;
    fopt = f(x);
    nopt = k;

end

function Z = f(x)
    Z = 0.5*x(1)^2 + 2.5*x(2)^2;
end

function Z = gradf(x)
    Z = [x(1);5*x(2)];
end
```

Summary

- Methods for optimization in 1D have “counterparts” in methods for the solution of nonlinear equations:

Golden Search → Bisection

Parabolic Interpolation → Regula Falsi

Newton ($f(x) = 0$) → Newton ($f'(x) = 0$)

and resemble many of their properties (linear/quadratic convergence etc.).

- Multidimensional optimization requires knowledge of gradients and sometimes Hessians, which generalize first and second order derivatives.
- Steepest descent moves in the direction of negative gradient - results in zig-zag moves (a method called conjugate-gradient fixes this problem).

Summary

- ▶ Multidimensional Newton's method generalizes Newton's method for optimization in 1D. It is fast, but requires significant work (deriving the Hessian, and solving a linear system).
- ▶ BFGS is an extremely popular secant-updating method which works with an approximate Hessian.